

Basic Instructions

1. Enter your Name and UID in the provided space.
2. Do the assignment in the notebook itself
3. you are free to use Google Colab

Name: **Arpit Aggarwal**

UID: **116747189**

In the first part, you will implement all the functions required to build a two layer neural network. In the next part, you will use these functions for image and text classification. Provide your code at the appropriate placeholders.

1. Packages

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import h5py
import scipy
from PIL import Image
from scipy import ndimage
```

2. Layer Initialization

Exercise: Create and initialize the parameters of the 2-layer neural network. Use random initialization for the weight matrices and zero initialization for the biases.

```
In [2]: def initialize_parameters(n_x, n_h, n_y):
        """
        Argument:
        n_x -- size of the input layer
        n_h -- size of the hidden layer
        n_y -- size of the output layer

        Returns:
        parameters -- python dictionary containing your parameters:
                        W1 -- weight matrix of shape (n_h, n_x)
                        b1 -- bias vector of shape (n_h, 1)
                        W2 -- weight matrix of shape (n_y, n_h)
                        b2 -- bias vector of shape (n_y, 1)

        """

        np.random.seed(1)

        ### START CODE HERE ### (~ 4 lines of code)

        W1 = np.random.randn(n_h, n_x) * 0.01
        b1 = np.zeros(shape=(n_h, 1))
        W2 = np.random.randn(n_y, n_h) * 0.01
        b2 = np.zeros(shape=(n_y, 1))

        ### END CODE HERE ###

        assert(W1.shape == (n_h, n_x))
        assert(b1.shape == (n_h, 1))
        assert(W2.shape == (n_y, n_h))
        assert(b2.shape == (n_y, 1))

        parameters = {"W1": W1,
                       "b1": b1,
                       "W2": W2,
                       "b2": b2}

        return parameters
```

```
In [3]: parameters = initialize_parameters(3,2,1)
        print("W1 = " + str(parameters["W1"]))
        print("b1 = " + str(parameters["b1"]))
        print("W2 = " + str(parameters["W2"]))
        print("b2 = " + str(parameters["b2"]))

        W1 = [[ 0.01624345 -0.00611756 -0.00528172]
               [-0.01072969  0.00865408 -0.02301539]]
        b1 = [[0.]
               [0.]]
        W2 = [[ 0.01744812 -0.00761207]]
        b2 = [[0.]]
```

Expected output:

```

**W1**      [[ 0.01624345 -0.00611756 -0.00528172]
              [-0.01072969 0.00865408 -0.02301539]]

**b1**      [[ 0.] [ 0.]]

**W2**      [[ 0.01744812 -0.00761207]]

**b2**      [[ 0.]]

```

3. Forward Propagation

Now that you have initialized your parameters, you will do the forward propagation module. You will start by implementing some basic functions that you will use later when implementing the model. You will complete three functions in this order:

- LINEAR
- LINEAR -> ACTIVATION where ACTIVATION will be either ReLU or Sigmoid.

The linear module computes the following equation:

$$Z = WA + b \quad (4)$$

3.1 Exercise: Build the linear part of forward propagation.

```
In [4]: def linear_forward(A, W, b):
        """
        Implement the linear part of a layer's forward propagation.

        Arguments:
        A -- activations from previous layer (or input data): (size of previous layer, number of examples)
        W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)
        b -- bias vector, numpy array of shape (size of the current layer, 1)

        Returns:
        Z -- the input of the activation function, also called pre-activation parameter
        cache -- a python dictionary containing "A", "W" and "b" ; stored for computing the backward pass efficiently
        """

        ### START CODE HERE ### (~ 1 line of code)

        Z = np.dot(W, A) + b

        ### END CODE HERE ###

        assert(Z.shape == (W.shape[0], A.shape[1]))
        cache = (A, W, b)

        return Z, cache
```

```
In [5]: np.random.seed(1)

A = np.random.randn(3,2)
W = np.random.randn(1,3)
b = np.random.randn(1,1)

Z, linear_cache = linear_forward(A, W, b)
print("Z = " + str(Z))
```

```
Z = [[ 3.26295337 -1.23429987]]
```

Expected output:

```

**Z**
[[ 3.26295337
 -1.23429987]]

```

3.2 - Linear-Activation Forward

In this notebook, you will use two activation functions:

- **Sigmoid:** $\sigma(Z) = \sigma(WA + b) = \frac{1}{1 + e^{-(WA+b)}}$. Write the code for the `sigmoid` function. This function returns **two** items: the activation value "a" and a "cache" that contains "Z" (it's what we will feed in to the corresponding backward function). To use it you could just call:

```
A, activation_cache = sigmoid(Z)
```

- **ReLU:** The mathematical formula for ReLU is $A = RELU(Z) = \max(0, Z)$. Write the code for the `relu` function. This function returns **two** items: the activation value "A" and a "cache" that contains "Z" (it's what we will feed in to the corresponding backward function). To use it you could just call: ``python A, activation_cache = relu(Z)

Exercise:

- Implement the activation functions
- Build the linear activation part of forward propagation. Mathematical relation is:

$$A = g(Z) = g(WA_{prev} + b)$$

```
In [6]: def sigmoid(Z):  
        """  
        Implements the sigmoid activation in numpy  
  
        Arguments:  
        Z -- numpy array of any shape  
  
        Returns:  
        A -- output of sigmoid(z), same shape as Z  
        cache -- returns Z, useful during backpropagation  
        """  
        ### START CODE HERE ### (~ 2 line of code)  
  
        A = 1.0 / (1.0 + np.exp(-Z))  
        cache = Z  
  
        ### END CODE HERE ###  
  
        return A, cache  
  
def relu(Z):  
    """  
    Implement the RELU function.  
  
    Arguments:  
    Z -- Output of the linear layer, of any shape  
  
    Returns:  
    A -- Post-activation parameter, of the same shape as Z  
    cache -- returns Z, useful during backpropagation  
    """  
    ### START CODE HERE ### (~ 2 line of code)  
  
    A = np.maximum(0, Z)  
    cache = Z  
  
    ### END CODE HERE ###  
  
    assert(A.shape == Z.shape)  
    return A, cache
```

```

In [7]: def linear_activation_forward(A_prev, W, b, activation):
        """
        Implement the forward propagation for the LINEAR->ACTIVATION layer

        Arguments:
        A_prev -- activations from previous layer (or input data): (size
        of previous layer, number of examples)
        W -- weights matrix: numpy array of shape (size of current layer,
        size of previous layer)
        b -- bias vector, numpy array of shape (size of the current layer,
        1)
        activation -- the activation to be used in this layer, stored as
        a text string: "sigmoid" or "relu"

        Returns:
        A -- the output of the activation function, also called the post-
        activation value
        cache -- a python dictionary containing "linear_cache" and "activation_cache";
        stored for computing the backward pass efficiently
        """

        if activation == "sigmoid":
            # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
            ### START CODE HERE ### (~ 2 lines of code)

            Z, linear_cache = linear_forward(A_prev, W, b)
            A, activation_cache = sigmoid(Z)

            ### END CODE HERE ###

        elif activation == "relu":
            # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
            ### START CODE HERE ### (~ 2 lines of code)

            Z, linear_cache = linear_forward(A_prev, W, b)
            A, activation_cache = relu(Z)

            ### END CODE HERE ###

        assert (A.shape == (W.shape[0], A_prev.shape[1]))
        cache = (linear_cache, activation_cache)

        return A, cache

```

```
In [8]: np.random.seed(2)
A_prev = np.random.randn(3,2)
W = np.random.randn(1,3)
b = np.random.randn(1,1)

A, linear_activation_cache = linear_activation_forward(A_prev, W, b,
activation = "sigmoid")
print("With sigmoid: A = " + str(A))

A, linear_activation_cache = linear_activation_forward(A_prev, W, b,
activation = "relu")
print("With ReLU: A = " + str(A))

With sigmoid: A = [[0.96890023 0.11013289]]
With ReLU: A = [[3.43896131 0.      ]]
```

Expected output:

```
**With sigmoid: A **      [[ 0.96890023
                             0.11013289]]

**With ReLU: A **      [[ 3.43896131 0.  ]]
```

4 - Loss function

Now you will implement forward and backward propagation. You need to compute the loss, because you want to check if your model is actually learning.

Exercise: Compute the cross-entropy loss J , using the following formula:

$$-\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{(i)}) + (1 - y^{(i)}) \log(1 - a^{(i)})) \quad (7)$$


```
In [9]: # GRADED FUNCTION: compute_loss

def compute_loss(A, Y):
    """
    Implement the loss function defined by equation (7).

    Arguments:
    A -- probability vector corresponding to your label predictions,
    shape (1, number of examples)
    Y -- true "label" vector (for example: containing 0 if non-cat, 1
    if cat), shape (1, number of examples)

    Returns:
    loss -- cross-entropy loss
    """

    m = Y.shape[1]

    # Compute loss from aL and y.
    ### START CODE HERE ### (~ 1 lines of code)

    loss = (-1.0 / m) * np.sum((Y * np.log(A)) + ((1.0 - Y) * np.log(
    1.0 - A)))

    ### END CODE HERE ###

    loss = np.squeeze(loss)      # To make sure your loss's shape is
    what we expect (e.g. this turns [[17]] into 17).
    assert(loss.shape == ())

    return loss
```

```
In [10]: Y = np.asarray([[1, 1, 1]])
A = np.array([[.8, .9, 0.4]])

print("loss = " + str(compute_loss(A, Y)))

loss = 0.41493159961539694
```

Expected Output:

```
**loss**  0.41493159961539694
```

5 - Backward propagation module

Just like with forward propagation, you will implement helper functions for backpropagation. Remember that back propagation is used to calculate the gradient of the loss function with respect to the parameters.

Now, similar to forward propagation, you are going to build the backward propagation in two steps:

- LINEAR backward
- LINEAR -> ACTIVATION backward where ACTIVATION computes the derivative of either the ReLU or sigmoid activation

5.1 - Linear backward

```
In [11]: # GRADED FUNCTION: linear_backward

def linear_backward(dZ, cache):
    """
    Implement the linear portion of backward propagation for a single
    layer (layer l)

    Arguments:
    dZ -- Gradient of the loss with respect to the linear output (of
    current layer l)
    cache -- tuple of values (A_prev, W, b) coming from the forward p
    ropagation in the current layer

    Returns:
    dA_prev -- Gradient of the loss with respect to the activation (o
    f the previous layer l-1), same shape as A_prev
    dW -- Gradient of the loss with respect to W (current layer l), s
    ame shape as W
    db -- Gradient of the loss with respect to b (current layer l), s
    ame shape as b
    """
    A_prev, W, b = cache
    m = A_prev.shape[1]

    ### START CODE HERE ### (~ 3 lines of code)

    dA_prev = np.dot(W.T, dZ)
    dW = np.dot(dZ, A_prev.T)
    db = np.array([np.sum(dZ, axis = 1)]).T

    ### END CODE HERE ###

    assert (dA_prev.shape == A_prev.shape)
    assert (dW.shape == W.shape)
    assert (db.shape == b.shape)

    return dA_prev, dW, db
```

```
In [12]: np.random.seed(1)
dZ = np.random.randn(1,2)
A = np.random.randn(3,2)
W = np.random.randn(1,3)
b = np.random.randn(1,1)
linear_cache = (A, W, b)

dA_prev, dW, db = linear_backward(dZ, linear_cache)
print ("dA_prev = " + str(dA_prev))
print ("dW = " + str(dW))
print ("db = " + str(db))

dA_prev = [[ 0.51822968 -0.19517421]
 [-0.40506361  0.15255393]
 [ 2.37496825 -0.89445391]]
dW = [[-0.2015379  2.81370193  3.2998501 ]]
db = [[1.01258895]]
```

Expected Output:

```
**dA_prev**  [[ 0.51822968 -0.19517421] [-0.40506361 0.15255393] [
               2.37496825 -0.89445391]]

**dW**              [[-0.2015379 2.81370193 3.2998501 ]]

**db**              [[1.01258895]]
```

5.2 - Linear Activation backward

Next, you will create a function that merges the two helper functions: **linear_backward** and the backward step for the activation **linear_activation_backward**.

Before implementing **linear_activation_backward**, you need to implement two backward functions for each activations:

- **sigmoid_backward** : Implements the backward propagation for SIGMOID unit. You can call it as follows:

```
dZ = sigmoid_backward(dA, activation_cache)
```

- **relu_backward** : Implements the backward propagation for RELU unit. You can call it as follows:

```
dZ = relu_backward(dA, activation_cache)
```

If $g(.)$ is the activation function, **sigmoid_backward** and **relu_backward** compute

$$dZ^{[l]} = dA^{[l]} * g'(Z^{[l]}) \quad (11)$$

Exercise:

- Implement the backward functions for the relu and sigmoid activation layer.
- Implement the backpropagation for the *LINEAR->ACTIVATION* layer.

```

In [13]: def relu_backward(dA, cache):
    """
    Implement the backward propagation for a single RELU unit.

    Arguments:
    dA -- post-activation gradient, of any shape
    cache -- 'Z' where we store for computing backward propagation efficiently

    Returns:
    dZ -- Gradient of the loss with respect to Z
    """

    Z = cache
    dZ = np.array(dA, copy=True) # just converting dz to a correct object.

    ### START CODE HERE ### (~ 1 line of code)

    dZ = dA * np.where(Z <= 0, 0, 1)

    ### END CODE HERE ###

    assert (dZ.shape == Z.shape)

    return dZ

def sigmoid_backward(dA, cache):
    """
    Implement the backward propagation for a single SIGMOID unit.

    Arguments:
    dA -- post-activation gradient, of any shape
    cache -- 'Z' where we store for computing backward propagation efficiently

    Returns:
    dZ -- Gradient of the loss with respect to Z
    """

    Z = cache

    ### START CODE HERE ### (~ 2 line of code)

    sigmoid_derivative = sigmoid(Z)[0] * (1.0 - sigmoid(Z)[0])
    dZ = dA * sigmoid_derivative

    ### END CODE HERE ###

    assert (dZ.shape == Z.shape)

    return dZ

```

```

In [14]: # GRADED FUNCTION: linear_activation_backward

def linear_activation_backward(dA, cache, activation):
    """
    Implement the backward propagation for the LINEAR->ACTIVATION layer.

    Arguments:
    dA -- post-activation gradient for current layer l
    cache -- tuple of values (linear_cache, activation_cache) we store
    for computing backward propagation efficiently
    activation -- the activation to be used in this layer, stored as
    a text string: "sigmoid" or "relu"

    Returns:
    dA_prev -- Gradient of the loss with respect to the activation (of
    the previous layer l-1), same shape as A_prev
    dW -- Gradient of the loss with respect to W (current layer l), same
    shape as W
    db -- Gradient of the loss with respect to b (current layer l), same
    shape as b
    """
    linear_cache, activation_cache = cache

    if activation == "relu":
        ### START CODE HERE ### (~ 2 lines of code)

        dZ = relu_backward(dA, activation_cache)
        dA_prev, dW, db = linear_backward(dZ, linear_cache)

        ### END CODE HERE ###

    elif activation == "sigmoid":
        ### START CODE HERE ### (~ 2 lines of code)

        dZ = sigmoid_backward(dA, activation_cache)
        dA_prev, dW, db = linear_backward(dZ, linear_cache)

        ### END CODE HERE ###

    return dA_prev, dW, db

```

```
In [15]: np.random.seed(2)
dA = np.random.randn(1,2)
A = np.random.randn(3,2)
W = np.random.randn(1,3)
b = np.random.randn(1,1)
Z = np.random.randn(1,2)
linear_cache = (A, W, b)
activation_cache = Z
linear_activation_cache = (linear_cache, activation_cache)

dA_prev, dW, db = linear_activation_backward(dA, linear_activation_cache, activation = "sigmoid")
print ("sigmoid:")
print ("dA_prev = " + str(dA_prev))
print ("dW = " + str(dW))
print ("db = " + str(db) + "\n")

dA_prev, dW, db = linear_activation_backward(dA, linear_activation_cache, activation = "relu")
print ("relu:")
print ("dA_prev = " + str(dA_prev))
print ("dW = " + str(dW))
print ("db = " + str(db))
```

```
sigmoid:
dA_prev = [[ 0.11017994  0.01105339]
 [ 0.09466817  0.00949723]
 [-0.05743092 -0.00576154]]
dW = [[ 0.20533573  0.19557101 -0.03936168]]
db = [[-0.11459244]]

relu:
dA_prev = [[ 0.44090989  0.          ]
 [ 0.37883606  0.          ]
 [-0.2298228  0.          ]]
dW = [[ 0.89027649  0.74742835 -0.20957978]]
db = [[-0.41675785]]
```

Expected output with sigmoid:

dA_prev	[[0.11017994 0.01105339] [0.09466817 0.00949723] [-0.05743092 -0.00576154]]
dW	[[0.20533573 0.19557101 -0.03936168]]
db	[[-0.11459244]]

Expected output with relu:

dA_prev	[[0.44090989 0.] [0.37883606 0.] [-0.2298228 0.]]
dW	[[0.89027649 0.74742835 -0.20957978]]
db	[[-0.41675785]]

6 - Update Parameters

In this section you will update the parameters of the model, using gradient descent:

$$W^{[1]} = W^{[1]} - \alpha dW^{[1]} \quad (16)$$

$$b^{[1]} = b^{[1]} - \alpha db^{[1]} \quad (17)$$

$$W^{[2]} = W^{[2]} - \alpha dW^{[2]} \quad (16)$$

$$b^{[2]} = b^{[2]} - \alpha db^{[2]} \quad (17)$$

where α is the learning rate. After computing the updated parameters, store them in the parameters dictionary.

Exercise: Implement `update_parameters()` to update your parameters using gradient descent.

Instructions: Update parameters using gradient descent.

```
In [16]: # GRADED FUNCTION: update_parameters

def update_parameters(parameters, grads, learning_rate):
    """
    Update parameters using gradient descent

    Arguments:
    parameters -- python dictionary containing your parameters
    grads -- python dictionary containing your gradients, output of L
              _model_backward

    Returns:
    parameters -- python dictionary containing your updated parameters

    """
    # Update rule for each parameter. Use a for loop.
    ### START CODE HERE ### (~ 4 lines of code)

    for key in parameters:
        parameters[key] = parameters[key] - (learning_rate * grads[
"d" + str(key)])

    ### END CODE HERE ###
    return parameters
```

```

In [17]: np.random.seed(2)
W1 = np.random.randn(3,4)
b1 = np.random.randn(3,1)
W2 = np.random.randn(1,3)
b2 = np.random.randn(1,1)
parameters = {"W1": W1,
              "b1": b1,
              "W2": W2,
              "b2": b2}

np.random.seed(3)
dW1 = np.random.randn(3,4)
db1 = np.random.randn(3,1)
dW2 = np.random.randn(1,3)
db2 = np.random.randn(1,1)
grads = {"dW1": dW1,
         "db1": db1,
         "dW2": dW2,
         "db2": db2}

parameters = update_parameters(parameters, grads, 0.1)

print ("W1 = "+ str(parameters["W1"]))
print ("b1 = "+ str(parameters["b1"]))
print ("W2 = "+ str(parameters["W2"]))
print ("b2 = "+ str(parameters["b2"]))

W1 = [[-0.59562069 -0.09991781 -2.14584584  1.82662008]
      [-1.76569676 -0.80627147  0.51115557 -1.18258802]
      [-1.0535704  -0.86128581  0.68284052  2.20374577]]
b1 = [[-0.04659241]
      [-1.28888275]
      [ 0.53405496]]
W2 = [[-0.55569196  0.0354055  1.32964895]]
b2 = [[-0.84610769]]

```

Expected Output:

W1	[[-0.59562069 -0.09991781 -2.14584584 1.82662008] [-1.76569676 -0.80627147 0.51115557 -1.18258802] [-1.0535704 -0.86128581 0.68284052 2.20374577]]
b1	[[-0.04659241] [-1.28888275] [0.53405496]]
W2	[[-0.55569196 0.0354055 1.32964895]]
b2	[[-0.84610769]]

7 - Conclusion

Congrats on implementing all the functions required for building a deep neural network!

We know it was a long assignment but going forward it will only get better. The next part of the assignment is easier.

Part 2:

In the next part you will put all these together to build a two-layer neural networks for image classification.

```
In [18]: %matplotlib inline
plt.rcParams['figure.figsize'] = (5.0, 4.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

np.random.seed(1)
```

Dataset

Problem Statement: You are given a dataset ("data/train_catvnoncat.h5", "data/test_catvnoncat.h5") containing:

- a training set of `m_train` images labelled as cat (1) or non-cat (0)
- a test set of `m_test` images labelled as cat and non-cat
- each image is of shape `(num_px, num_px, 3)` where 3 is for the 3 channels (RGB).

Let's get more familiar with the dataset. Load the data by completing the function and run the cell below.

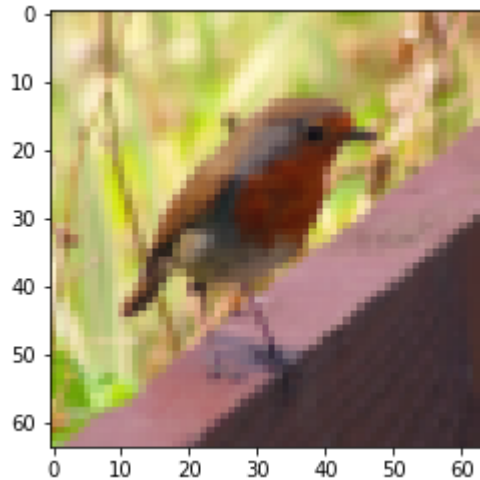
```
In [19]: def load_data(train_file, test_file):  
    # Load the training data  
    train_dataset = h5py.File(train_file, 'r')  
  
    # Separate features(x) and labels(y) for training set  
    train_set_x_orig = np.array(train_dataset['train_set_x'])  
    train_set_y_orig = np.array(train_dataset['train_set_y'])  
  
    # Load the test data  
    test_dataset = h5py.File(test_file, 'r')  
  
    # Separate features(x) and labels(y) for training set  
    test_set_x_orig = np.array(test_dataset['test_set_x'])  
    test_set_y_orig = np.array(test_dataset['test_set_y'])  
  
    classes = np.array(test_dataset["list_classes"][:]) # the list of  
    classes  
  
    train_set_y_orig = train_set_y_orig.reshape((1, train_set_y_orig.  
    shape[0]))  
    test_set_y_orig = test_set_y_orig.reshape((1, test_set_y_orig.sha  
    pe[0]))  
  
    return train_set_x_orig, train_set_y_orig, test_set_x_orig, test_  
    set_y_orig, classes
```

```
In [20]: train_file="data/train_catvnoncat.h5"  
    test_file="data/test_catvnoncat.h5"  
    train_x_orig, train_y, test_x_orig, test_y, classes = load_data(train_  
    _file, test_file)
```

The following code will show you an image in the dataset. Feel free to change the index and re-run the cell multiple times to see other images.

```
In [21]: # Example of a picture
index = 10
plt.imshow(train_x_orig[index])
print ("y = " + str(train_y[0,index]) + ". It's a " + classes[train_y
[0,index]].decode("utf-8") + " picture.")
```

y = 0. It's a non-cat picture.



```
In [22]: # Explore your dataset
m_train = train_x_orig.shape[0]
num_px = train_x_orig.shape[1]
m_test = test_x_orig.shape[0]

print ("Number of training examples: " + str(m_train))
print ("Number of testing examples: " + str(m_test))
print ("Each image is of size: (" + str(num_px) + ", " + str(num_px)
+ ", 3)")
print ("train_x_orig shape: " + str(train_x_orig.shape))
print ("train_y shape: " + str(train_y.shape))
print ("test_x_orig shape: " + str(test_x_orig.shape))
print ("test_y shape: " + str(test_y.shape))
```

Number of training examples: 209
 Number of testing examples: 50
 Each image is of size: (64, 64, 3)
 train_x_orig shape: (209, 64, 64, 3)
 train_y shape: (1, 209)
 test_x_orig shape: (50, 64, 64, 3)
 test_y shape: (1, 50)

As usual, you reshape and standardize the images before feeding them to the network.



Figure 1: Image to vector conversion.

```
In [23]: # Reshape the training and test examples
train_x_flatten = train_x_orig.reshape(train_x_orig.shape[0], -1).T
# The "-1" makes reshape flatten the remaining dimensions
test_x_flatten = test_x_orig.reshape(test_x_orig.shape[0], -1).T

# Standardize data to have feature values between 0 and 1.
train_x = train_x_flatten/255.
test_x = test_x_flatten/255.

print ("train_x's shape: " + str(train_x.shape))
print ("test_x's shape: " + str(test_x.shape))

train_x's shape: (12288, 209)
test_x's shape: (12288, 50)
```

3 - Architecture of your model

Now that you are familiar with the dataset, it is time to build a deep neural network to distinguish cat images from non-cat images.

2-layer neural network



Figure 2: 2-layer neural network.

The model can be summarized as: *****INPUT -> LINEAR -> RELU -> LINEAR -> SIGMOID -> OUTPUT*****.

Detailed Architecture of figure 2:

- The input is a (64,64,3) image which is flattened to a vector of size (12288, 1).
- The corresponding vector: $[x_0, x_1, \dots, x_{12287}]^T$ is then multiplied by the weight matrix $W^{[1]}$ of size $(n^{[1]}, 12288)$.
- You then add a bias term and take its relu to get the following vector: $[a_0^{[1]}, a_1^{[1]}, \dots, a_{n^{[1]}-1}^{[1]}]^T$.
- You multiply the resulting vector by $W^{[2]}$ and add your intercept (bias).
- Finally, you take the sigmoid of the result. If it is greater than 0.5, you classify it to be a cat.

General methodology

As usual you will follow the Deep Learning methodology to build the model:

1. Initialize parameters / Define hyperparameters
2. Loop for num_iterations:
 - a. Forward propagation
 - b. Compute loss function
 - c. Backward propagation
 - d. Update parameters (using parameters, and grads from backprop)

Question: Use the helper functions you have implemented in the previous assignment to build a 2-layer neural network with the following structure: *LINEAR* -> *RELU* -> *LINEAR* -> *SIGMOID*. The functions you may need and their inputs are:

```
def initialize_parameters(n_x, n_h, n_y):
    ...
    return parameters
def linear_activation_forward(A_prev, W, b, activation):
    ...
    return A, cache
def compute_loss(AL, Y):
    ...
    return loss
def linear_activation_backward(dA, cache, activation):
    ...
    return dA_prev, dW, db
def update_parameters(parameters, grads, learning_rate):
    ...
    return parameters
```

```
In [113]: ### CONSTANTS DEFINING THE MODEL ###
n_x = 12288      # num_px * num_px * 3
n_h = 15
n_y = 1
layers_dims = (n_x, n_h, n_y)
```

```

In [114]: def two_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_ite
          erations = 3000, print_loss=False):
          """
          Implements a two-layer neural network: LINEAR->RELU->LINEAR->SIGM
          OID.

          Arguments:
          X -- input data, of shape (n_x, number of examples)
          Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of
          shape (1, number of examples)
          layers_dims -- dimensions of the layers (n_x, n_h, n_y)
          num_iterations -- number of iterations of the optimization loop
          learning_rate -- learning rate of the gradient descent update rul
          e
          print_loss -- If set to True, this will print the loss every 100
          iterations

          Returns:
          parameters -- a dictionary containing W1, W2, b1, and b2
          """

          np.random.seed(1)
          grads = {}
          losses = []                                     # to keep track of the l
          oss
          m = X.shape[1]                                  # number of examples
          (n_x, n_h, n_y) = layers_dims

          # Initialize parameters dictionary, by calling one of the functio
          ns you'd previously implemented
          ### START CODE HERE ### (~ 1 line of code)

          parameters = initialize_parameters(n_x, n_h, n_y)

          ### END CODE HERE ###

          # Get W1, b1, W2 and b2 from the dictionary parameters.
          W1 = parameters["W1"]
          b1 = parameters["b1"]
          W2 = parameters["W2"]
          b2 = parameters["b2"]

          # Loop (gradient descent)

          for i in range(0, num_iterations):

              # Forward propagation: LINEAR -> RELU -> LINEAR -> SIGMOID. I
              nputs: "X, W1, b1, W2, b2". Output: "A1, cache1, A2, cache2".
              ### START CODE HERE ### (~ 2 lines of code)

              A1, cache1 = linear_activation_forward(X, W1, b1, "relu")
              A2, cache2 = linear_activation_forward(A1, W2, b2, "sigmoid")

              ### END CODE HERE ###

              # Compute loss

```

```

    ### START CODE HERE ### (~ 1 line of code)

    loss = compute_loss(A2, Y)

    ### END CODE HERE ###

    # Initializing backward propagation
    dA2 = - (np.divide(Y, A2) - np.divide(1 - Y, 1 - A2))/m

    # Backward propagation. Inputs: "dA2, cache2, cache1". Output
s: "dA1, dW2, db2; also dA0 (not used), dW1, db1".
    ### START CODE HERE ### (~ 2 lines of code)

    dA1, dW2, db2 = linear_activation_backward(dA2, cache2, "sigm
oid")
    dA0, dW1, db1 = linear_activation_backward(dA1, cache1, "rel
u")

    ### END CODE HERE ###

    # Set grads['dW1'] to dW1, grads['db1'] to db1, grads['dW2']
to dW2, grads['db2'] to db2
    ### START CODE HERE ### (~ 4 lines of code)

    grads['dW1'] = dW1
    grads['db1'] = db1
    grads['dW2'] = dW2
    grads['db2'] = db2

    ### END CODE HERE ###

    # Update parameters.
    ### START CODE HERE ### (approx. 1 line of code)

    parameters = update_parameters(parameters, grads, learning_ra
te)

    ### END CODE HERE ###

    # Retrieve W1, b1, W2, b2 from parameters
    W1 = parameters["W1"]
    b1 = parameters["b1"]
    W2 = parameters["W2"]
    b2 = parameters["b2"]

    # Print the loss every 100 training example
    if print_loss and i % 100 == 0:
        print("Loss after iteration {}: {}".format(i, np.squeeze(
loss)))
    if print_loss and i % 100 == 0:
        losses.append(loss)

    # plot the loss

    plt.plot(np.squeeze(losses))
    plt.ylabel('loss')

```




```
plt.xlabel('iterations (per tens)')
plt.title("Learning rate =" + str(learning_rate))
plt.show()

return parameters
```

In [130]: parameters = two_layer_model(train_x, train_y, layers_dims = (n_x, n_h, n_y), learning_rate=0.005, num_iterations = 3000, print_loss=True)

```
Loss after iteration 0: 0.695340595645
Loss after iteration 100: 0.649548208986
Loss after iteration 200: 0.642676975158
Loss after iteration 300: 0.631660284394
Loss after iteration 400: 0.611600199222
Loss after iteration 500: 0.584694561332
Loss after iteration 600: 0.554108895641
Loss after iteration 700: 0.519408709446
Loss after iteration 800: 0.479346954263
Loss after iteration 900: 0.435289848899
Loss after iteration 1000: 0.389759684786
Loss after iteration 1100: 0.345536220696
Loss after iteration 1200: 0.304332345327
Loss after iteration 1300: 0.26739268936
Loss after iteration 1400: 0.235050695399
Loss after iteration 1500: 0.207926871865
Loss after iteration 1600: 0.183558187172
Loss after iteration 1700: 0.157476388129
Loss after iteration 1800: 0.13816776328
Loss after iteration 1900: 0.123984741236
Loss after iteration 2000: 0.10718025021
Loss after iteration 2100: 0.096021444701
Loss after iteration 2200: 0.0846817819257
Loss after iteration 2300: 0.0759370533833
Loss after iteration 2400: 0.0683579103607
Loss after iteration 2500: 0.0618350235786
Loss after iteration 2600: 0.0562150908573
Loss after iteration 2700: 0.0513348887243
Loss after iteration 2800: 0.0470535200906
Loss after iteration 2900: 0.0432182189221
```



Expected Output:

```
      **Loss after iteration 0**      0.6930497356599888
      **Loss after iteration 100**    0.6464320953428849
                                ** ... **
      **Loss after iteration 2400**    0.048554785628770206
```

Good thing you built a vectorized implementation! Otherwise it might have taken 10 times longer to train this.

Now, you can use the trained parameters to classify images from the dataset.

Exercise:

- Implement the forward function
- Implement the predict function below to make prediction on test_images

```

In [131]: def two_layer_forward(X, parameters):
          """
          Implement forward propagation for the [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID computation

          Arguments:
          X -- data, numpy array of shape (input size, number of examples)
          parameters -- output of initialize_parameters_deep()

          Returns:
          A2 -- last post-activation value
          caches -- list of caches containing:
                        every cache of linear_relu_forward() (there are L-1 of them, indexed from 0 to L-2)
                        the cache of linear_sigmoid_forward() (there is one, indexed L-1)
          """

          caches = []
          A = X

          # Implement LINEAR -> RELU. Add "cache" to the "caches" list.
          ### START CODE HERE ### (approx. 3 line of code)

          W1, b1 = parameters["W1"], parameters["b1"]
          A1, cache1 = linear_activation_forward(A, W1, b1, "relu")
          caches.append(cache1)

          ### END CODE HERE ###

          # Implement LINEAR -> SIGMOID. Add "cache" to the "caches" list.
          ### START CODE HERE ### (approx. 3 line of code)

          W2, b2 = parameters["W2"], parameters["b2"]
          A2, cache2 = linear_activation_forward(A1, W2, b2, "sigmoid")
          caches.append(cache2)

          ### END CODE HERE ###

          assert(A2.shape == (1,X.shape[1]))

          return A2, caches

```

```
In [132]: def predict(X, y, parameters):
    """
    This function is used to predict the results of a L-layer neural
    network.

    Arguments:
    X -- data set of examples you would like to label
    parameters -- parameters of the trained model

    Returns:
    p -- predictions for the given dataset X
    """

    m = X.shape[1]
    n = len(parameters) // 2 # number of layers in the neural network
    p = np.zeros((1,m))

    # Forward propagation
    ### START CODE HERE ### (~ 1 lines of code)

    probas, caches = two_layer_forward(X, parameters)

    ### END CODE HERE ###

    # convert probas to 0/1 predictions
    for i in range(0, probas.shape[1]):
        ### START CODE HERE ### (~ 4 lines of code)

        if(probas[0][i] > 0.5):
            p[0][i] = 1
        else:
            p[0][i] = 0

        ### END CODE HERE ###

    print("Accuracy: " + str(float(np.sum((p == y)/float(m)))))
    return p
```

```
In [133]: predictions_train = predict(train_x, train_y, parameters)
```

Accuracy: 1.0

```
In [134]: predictions_test = predict(test_x, test_y, parameters)
```

Accuracy: 0.72

Exercise: Identify the hyperparameters in the model and For each hyperparameter

- Briefly explain its role
- Explore a range of values and describe their impact on (a) training loss and (b) test accuracy
- Report the best hyperparameter value found.

Note: Provide your results and explanations in the report for this question.

Hyperparameters The hyperparameters are:

1. Learning rate - It is used for updating the parameters of the neural network that is the weights and biases of the neural network. It controls the amount of update that needs to take place so that we are able to reach the minima of the loss function.
2. Epochs - It represents the number of times the network sees the data and adjusts its parameters for optimal learning.
3. Number of hidden neurons in the hidden layer - The number of neurons in the hidden layer where each neuron is learning some properties of the input data and able to establish a relationship between input and output.

Values of Hyperparameters tried:

1. Learning rate = 0.05, Epochs = 3000, Hidden neurons = 50, Training loss = 0.05, Testing accuracy: 74%
2. Learning rate = 0.01, Epochs = 3000, Hidden neurons = 50, Training loss = 0.05, Testing accuracy: 74%
3. Learning rate = 0.01, Epochs = 4000, Hidden neurons = 50, Training loss = 0.05, Testing accuracy: 74%
4. Learning rate = 0.001, Epochs = 3000, Hidden neurons = 50, Training loss = 0.52, Testing accuracy: 58%
5. Learning rate = 0.005, Epochs = 3000, Hidden neurons = 50, Training loss = 0.03, Testing accuracy: 70%
6. Learning rate = 0.05, Epochs = 3000, Hidden neurons = 40, Training loss = 0.001, Testing accuracy: 76%
7. Learning rate = 0.01, Epochs = 3000, Hidden neurons = 40, Training loss = 0.01, Testing accuracy: 72%
8. Learning rate = 0.01, Epochs = 4000, Hidden neurons = 40, Training loss = 0.02, Testing accuracy: 72%
9. Learning rate = 0.001, Epochs = 3000, Hidden neurons = 40, Training loss = 0.52, Testing accuracy: 54%
10. Learning rate = 0.005, Epochs = 3000, Hidden neurons = 40, Training loss = 0.04, Testing accuracy: 74%
11. Learning rate = 0.05, Epochs = 3000, Hidden neurons = 30, Training loss = 0.005, Testing accuracy: 74%
12. Learning rate = 0.01, Epochs = 3000, Hidden neurons = 30, Training loss = 0.0011, Testing accuracy: 76%
13. Learning rate = 0.01, Epochs = 4000, Hidden neurons = 30, Training loss = 0.001, Testing accuracy: 76%
14. Learning rate = 0.001, Epochs = 3000, Hidden neurons = 30, Training loss = 0.53, Testing accuracy: 50%
15. Learning rate = 0.005, Epochs = 3000, Hidden neurons = 30, Training loss = 0.04, Testing accuracy: 72%
16. Learning rate = 0.05, Epochs = 3000, Hidden neurons = 20, Training loss = 0.0009, Testing accuracy: 78%
17. Learning rate = 0.01, Epochs = 3000, Hidden neurons = 20, Training loss = 0.011, Testing accuracy: 70%
18. Learning rate = 0.01, Epochs = 4000, Hidden neurons = 20, Training loss = 0.02, Testing accuracy: 70%
19. Learning rate = 0.001, Epochs = 3000, Hidden neurons = 20, Training loss = 0.54, Testing accuracy: 46%
20. Learning rate = 0.005, Epochs = 3000, Hidden neurons = 20, Training loss = 0.04, Testing accuracy: 72%
21. Learning rate = 0.05, Epochs = 3000, Hidden neurons = 15, Training loss = 0.005, Testing accuracy: 70%
22. Learning rate = 0.01, Epochs = 3000, Hidden neurons = 15, Training loss = 0.011, Testing accuracy: 74%
23. Learning rate = 0.01, Epochs = 4000, Hidden neurons = 15, Training loss = 0.02, Testing accuracy: 74%
24. Learning rate = 0.001, Epochs = 3000, Hidden neurons = 15, Training loss = 0.56, Testing accuracy: 36%
25. Learning rate = 0.005, Epochs = 3000, Hidden neurons = 15, Training loss = 0.043, Testing accuracy: 72%

Optimal hyperparameters found

1. Learning rate = 0.05
2. Epochs = 3000
3. Number of hidden neurons = 20

Results Analysis

First, let's take a look at some images the 2-layer model labeled incorrectly. This will show a few mislabeled images.

```
In [135]: def print_mislabeled_images(classes, X, y, p):
          """
          Plots images where predictions and truth were different.
          X -- dataset
          y -- true labels
          p -- predictions
          """
          a = p + y
          mislabeled_indices = np.asarray(np.where(a == 1))
          plt.rcParams['figure.figsize'] = (40.0, 40.0) # set default size
          of plots
          num_images = len(mislabeled_indices[0])
          for i in range(num_images):
              index = mislabeled_indices[1][i]

              plt.subplot(2, num_images, i + 1)
              plt.imshow(X[:,index].reshape(64,64,3), interpolation='nearest')
              plt.axis('off')
              plt.title("Prediction: " + classes[int(p[0,index])].decode("utf-8") + " \n Class: " + classes[y[0,index]].decode("utf-8"))
```

```
In [136]: print_mislabeled_images(classes, test_x, test_y, predictions_test)
```



Exercise: Identify a few types of images that tends to perform poorly on the model

Answer The model performs poorly when the cat is at certain angle or rotated at some angle, which makes it classify it as a non-cat class.

Now, lets use the same architecture to predict sentiment of movie reviews. In this section, most of the implementation is already provided. The exercises are mainly to understand what the workflow is when handling the text data.

```
In [137]: import re
```

Dataset

Problem Statement: You are given a dataset ("train_imdb.txt", "test_imdb.txt") containing:

- a training set of `m_train` reviews
- a test set of `m_test` reviews
- the labels for the training examples are such that the first 50% belong to class 1 (positive) and the rest 50% of the data belong to class 0 (negative)

Let's get more familiar with the dataset. Load the data by completing the function and run the cell below.

```
In [138]: def load_data(train_file, test_file):
            train_dataset = []
            test_dataset = []

            # Read the training dataset file line by line
            for line in open(train_file, 'r'):
                train_dataset.append(line.strip())

            for line in open(test_file, 'r'):
                test_dataset.append(line.strip())
            return train_dataset, test_dataset
```

```
In [139]: train_file = "data/train_imdb.txt"
            test_file = "data/test_imdb.txt"
            train_dataset, test_dataset = load_data(train_file, test_file)
```

```
In [140]: # This is just how the data is organized. The first 50% data is positive
            # and the rest 50% is negative for both train and test splits.
            y = [1 if i < len(train_dataset)*0.5 else 0 for i in range(len(train_
            dataset))]
```

As usual, let's check our dataset

```
In [141]: # Example of a review
            index = 10
            print(train_dataset[index])
            print("y = " + str(y[index]))
```

I liked the film. Some of the action scenes were very interesting, tense and well done. I especially liked the opening scene which had a semi truck in it. A very tense action scene that seemed well done.

Some of the transitional scenes were filmed in interesting ways such as time lapse photography, unusual colors, or interesting angles. Also the film is funny in several parts. I also liked how the evil guy was portrayed too. I'd give the film an 8 out of 10.
y = 1


```
In [142]: # Explore your dataset
m_train = len(train_dataset)
m_test = len(test_dataset)

print ("Number of training examples: " + str(m_train))
print ("Number of testing examples: " + str(m_test))
```

Number of training examples: 1001

Number of testing examples: 201

Pre-Processing

From the example review, you can see that the raw data is really noisy! This is generally the case with the text data. Hence, Preprocessing the raw input and cleaning the text is essential. Please run the code snippet provided below.

Exercise: Explain what pattern the model is trying to capture using re.compile in your report.

Answer

1. re.compile() removes special characters like ', . ' etc and makes all characters in lowercase. It is learning properties from words.

```
In [143]: REPLACE_NO_SPACE = re.compile("(\\.|\\;|\\:|\\!|\\'|\\?|\\,|\\(|\\)|\\[|\\]|\\d+)"
REPLACE_WITH_SPACE = re.compile("<br\\s*><br\\s*>|\\-|\\/")
NO_SPACE = ""
SPACE = " "

def preprocess_reviews(reviews):
    reviews = [REPLACE_NO_SPACE.sub(NO_SPACE, line.lower()) for line
in reviews]
    reviews = [REPLACE_WITH_SPACE.sub(SPACE, line) for line in review
s]

    return reviews

train_dataset_clean = preprocess_reviews(train_dataset)
test_dataset_clean = preprocess_reviews(test_dataset)
```

```
In [144]: # Example of a clean review
index = 10
print(train_dataset_clean[index])
print ("y = " + str(y[index]))
```

```
i liked the film some of the action scenes were very interesting tens
e and well done i especially liked the opening scene which had a semi
truck in it a very tense action scene that seemed well done some of t
he transitional scenes were filmed in interesting ways such as time l
apse photography unusual colors or interesting angles also the film i
s funny is several parts i also liked how the evil guy was portrayed
too id give the film an out of
y = 1
```

Vectorization

Now lets create a feature vector for our reviews based on a simple bag of words model. So, given an input text, we need to create a numerical vector which is simply the vector of word counts for each word of the vocabulary. Run the code below to get the feature representation.

```
In [145]: from sklearn.feature_extraction.text import CountVectorizer

cv = CountVectorizer(binary=True, stop_words="english", max_features=
2000)
cv.fit(train_dataset_clean)
X = cv.transform(train_dataset_clean)
X_test = cv.transform(test_dataset_clean)
```

CountVectorizer provides a sparse feature representation by default which is reasonable because only some words occur in individual example. However, for training neural network models, we generally use a dense representation vector.

```
In [146]: X = np.array(X.todense()).astype(float)
X_test = np.array(X_test.todense()).astype(float)
y = np.array(y)
```

Model

```
In [147]: from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split

X_train, X_val, y_train, y_val = train_test_split(
    X, y, train_size = 0.80
)
```

/home/arpitdec5/.local/lib/python2.7/site-packages/sklearn/model_selection/_split.py:2178: FutureWarning: From version 0.21, test_size will always complement train_size unless both are specified.
FutureWarning)


```
In [148]: # This is just to correct the shape of the arrays as required by the
two_layer_model
X_train = X_train.T
X_val = X_val.T
y_train = y_train.reshape(1,-1)
y_val = y_val.reshape(1,-1)
```

```
In [165]: ### CONSTANTS DEFINING THE MODEL ###
n_x = X_train.shape[0]
n_h = 800
n_y = 1
layers_dims = (n_x, n_h, n_y)
```

We will use the same two layer model that you completed in the previous section for training.

```
In [166]: parameters = two_layer_model(X_train, y_train, layers_dims = (n_x, n_h, n_y), learning_rate=0.05, num_iterations = 3000, print_loss=True)
```

```
Loss after iteration 0: 0.693741186555
Loss after iteration 100: 0.669243147014
Loss after iteration 200: 0.570034533211
Loss after iteration 300: 0.367406293499
Loss after iteration 400: 0.230506374002
Loss after iteration 500: 0.154438318239
Loss after iteration 600: 0.108656728844
Loss after iteration 700: 0.0794306058322
Loss after iteration 800: 0.0600857685661
Loss after iteration 900: 0.0469245239498
Loss after iteration 1000: 0.0376980944249
Loss after iteration 1100: 0.0310273960424
Loss after iteration 1200: 0.0260620102052
Loss after iteration 1300: 0.0222701491873
Loss after iteration 1400: 0.019308336846
Loss after iteration 1500: 0.016948431047
Loss after iteration 1600: 0.0150356278223
Loss after iteration 1700: 0.0134616437642
Loss after iteration 1800: 0.0121489007512
Loss after iteration 1900: 0.0110410883932
Loss after iteration 2000: 0.0100963185528
Loss after iteration 2100: 0.00928302482988
Loss after iteration 2200: 0.00857702273427
Loss after iteration 2300: 0.00795952326132
Loss after iteration 2400: 0.00741573609318
Loss after iteration 2500: 0.00693385723277
Loss after iteration 2600: 0.00650442367197
Loss after iteration 2700: 0.00611972158743
Loss after iteration 2800: 0.00577344849862
Loss after iteration 2900: 0.0054603771549
```



Predict the review for our movies!

```
In [167]: predictions_train = predict(X_train, y_train, parameters)
```

Accuracy: 1.0

```
In [168]: predictions_val = predict(X_val, y_val, parameters)
```

Accuracy: 0.8407960199

Results Analysis

Let's take a look at some examples the 2-layer model labeled incorrectly

```
In [169]: def print_mislabeled_reviews(X, y, p):  
    """  
    Plots images where predictions and truth were different.  
    X -- dataset  
    y -- true labels  
    p -- predictions  
    """  
  
    a = p + y  
    mislabeled_indices = np.asarray(np.where(a == 1))  
    plt.rcParams['figure.figsize'] = (40.0, 40.0) # set default size  
of plots  
    num_reviews = len(mislabeled_indices[0])  
    for i in range(num_reviews):  
        index = mislabeled_indices[1][i]  
  
        print((" ").join(cv.inverse_transform(X[index])[0]))  
        print("Prediction: " + str(int(p[0,index])) + " \n Class: " +  
str(y[0,index]))
```

```
In [170]: print_mislabeled_reviews(X_val.T, y_val, predictions_val)
```


achieve acting actor actors adaptation animal annoying author away background beautiful better book books camp care characters charismatic chosen cinematography clear complaint complex country critics despite different director doing drinking dull episode episodes especially exist fact feels film films forget friends good guess hand having heavy humour just know leads life like little lot main mainly make makes meet men moment movie need novel obvious opinion opposite party people person pleasant plot positive predictable presented presents psychological quiet quite rarely read reading real really recent relationship scenes screen screenplay script somewhat spends story strong suffers surprise surprised surprisingly taking talented thats thriller tim time totally twice ultimately van woman works writers writing wrote year

Prediction: 0

Class: 1

actually apparent brother close connected cop drive extreme far figure films forget gang getting great hard harsh involved kid killer lawyer look lot movie movies need ones place pretty scenes seen sex society sort start style takes themes tough trilogy true trying typical usual violent way ways worse worth you'll you've

Prediction: 0

Class: 1

animal ask aspect bad chair charlie check classic country desperate dont eventually fan film films final finale foreign forget friends gets gore greatest guts heard hell heres highly horror hours japanese job just kill kinds knowing like listen make movie people probably pure rating read really recommend said seen series shes shock sick snuff sound spin star storyline stuff things think throw times torture underground use watch watching woman

Prediction: 0

Class: 1

accident acting actually anti away bad beginning best better bit bring car character comes completely confusing create created decent device dialogue didnt die director effect elements episode expecting extreme face far feel felt film forth fresh future genres ghost gives goes going good gore grace great harder hero hes hope horrible horror idea images issue john key later left let like little live look looks main make man master nice opinion particularly plot prior protagonist read relationship remind revenge review reviewer revolves right said saving say season second seen sense solid somewhat spectacular state story strong suspense takes talking think time times turn twists viewer wants watch watching way wife wish wrong

Prediction: 0

Class: 1

accurate actors actual actually amazing bad best better big bother break budget came cameron capture care certainly cover critics day decent dialogue dicaprio did didnt dont effects enjoyed face film films fine form general going good guess guys heard hearing hit huge including isnt ive james job just lets like love make maybe million mind money mouth movie needed obviously people personally public real really romance say screen seeing seen ship sinking special star started story success sucks sweet talk talking thank theaters thing think thought time times titanic took touch twice type wasnt winslet word work worst worth wouldnt writers years

Prediction: 0

Class: 1

angry attempts author band based bit case characters clever compariso

n didnt directors disappointing documentary dream excellent family festival fiction film form genres good half just la lose lost main meant moving order original plus pretty previous quickly reminded rock roll saw science sequences seriously songs story sure sympathetic tell times tom way werent youll

Prediction: 1

Class: 0

animal dont hours life love planet script series video wait

Prediction: 1

Class: 0

actors art ask attempts audiences awful beautiful case character characters cinema cute date deliver despite development doesnt exactly expectations expected expecting experience fails female film flaws frank good hes high humour interesting know lead leads learn like liked main male matters men movie nice painful people plays potential pretty questions reviews romance romantic school screenplay shes simply story takes thats theyre tries try victim viewer women written youre

Prediction: 1

Class: 0

actual actually better bore bored called course easily ghost good human ive just like makes mindless movie original parts reporter revealed said slow st started story thats think thrown time town usually watch watched wow yes

Prediction: 1

Class: 0

acting actually added ask away basic biggest blame character characters children come coming convincing couple decent development didnt disappointing does doesnt dont english european experience female fully good great job king know later learns little lot lover main meant movie musical ok opinion piece played portrayal possibly presence really romance screenplay situation situations songs st stand story storyline things times understand viewer watched wife woman

Prediction: 1

Class: 0

actually aspect bad bit certainly chinese decent english explained extent family fighting film follow half hard head japanese just know language little make maybe middle plenty plus problem problems reading ride sexual simply storyline tried twists unlikely violence watching window

Prediction: 0

Class: 1

ability able accident action actresses actually aspect away bad believe better bit blood brothers cause cgi charlie crap crime cut days deal death disturbing does doesnt dont effects especially eyes fake favorite film films footage forget funny happens hope horror im instead leaving like look lot make makers making marry money movie movies overall people plot point porn probably pull rape rating real saying says scene scenes seen series shocking snuff sound stand stars sucked sucks super supposed sure talent talking thing thinking time tried visual visuals want wanted wasnt watch wouldnt

Prediction: 0

Class: 1

book character charlotte excellent eyes eyre goes hurt ice jane laid later like love mr novel passion read rochester saw story time totally version watching william wonderful years

Prediction: 1

Class: 0

actors ago believe better bit characters contains crazy dvd ed felt german good hard havent head high level like live look lot maybe missing movie movies nasty really saw say scenes seen short sick story thats thing uk version violence years

Prediction: 0

Class: 1

actually ahead better box budget burning character charlie come damn development did didnt dont enjoy eye fake far film flick forced genre good guinea guts hand hear heard horror hours interested just know like listen looks lot low making men minutes movie naturally offer painful pretty really recommend say scene scenes second seen set sharp short simply snuff story think thought throwing told torture trying ultimately unless various watching ways went woman worst

Prediction: 0

Class: 1

apparent audience better book boy certain certainly clear complete dark doesnt effort feels felt film good holes inspired intriguing like makes mystery nature portrayal premise probably production recommend tension thing true truly williams woman written

Prediction: 0

Class: 1

action age body brain building certainly computer crazy damme daughter dead entertaining especially fan fi fights folks genius goes going goldberg good government guess hes humor just keeps king lame later latest like manages mean named new original particularly perfect power pretty pro reason run sci sequel shoot site snake soldiers sort step super takes thriller train usual van war white working wrong year years youre

Prediction: 1

Class: 0

actors big deserves doesnt film goes great innocence movie real story tell think time victor

Prediction: 1

Class: 0

alive away basic brother camera characters cinema comes cool cop dead director doesnt ending ends expect experience fairly fallen film gang hard heart hollywood horrific humour incredibly isnt john just lawyer list look making moments movie nasty new nice old pace painful quite reveals said screen script seen sense showing stuff sudden takes taking think thriller time todays truly twice unexpected unusual utterly violence want way work working youre

Prediction: 0

Class: 1

absolutely actor art bad bed boss brings called came captured character chinese days definitely effort end english fact felt finish gem gets girl going great harder immediately just kept kid king little live love masks maybe movie nice night original pass perfectly really scott sees started thought totally touches trying turned turns watching

Prediction: 0

Class: 1

actually ago bad better book church course does enjoyable familiar film forgotten forward good hadnt heard hour instantly job know laid let long minute minutes missed mr nearly overall quick read really school second sense short simply story tales thats thing time trilogy watched worked write years

Prediction: 0

Class: 1

acting actor actors audience authentic belongs better changes charact
ers considered daughter director directors drinking endless essential
feelings fights films foreign good half head india indian involving l
eading lets like live make male movie movies nana old ones patekar re
d return running sad said scene school sell song state storytelling t
ale tell thrown tight tough true village violence violent walk wants
watch wont word words

Prediction: 0

Class: 1

acting actually appeal aspect audience bad ben big bit break brings c
aused completely consider damn decided did didnt died disbelief dolla
rs dont doubt dumb emotional emotions end enjoy epic exactly films fu
n girl gone good got greatest gross hate heavy idea im impact impress
ion isnt ive just kinda left let like little look make maybe means mi
llion mind mom movie movies order oscars possible power pretty quite
ready reason remember remembered ride saw score screen simply single
special sure theaters thoroughly thrill time times titanic total tv t
wice usually video visual wasnt whats whatsoever works worth years ye
s youll

Prediction: 0

Class: 1

anybody away ball believe big bit blown boat bodies bother car cares
cause chases course dead death did does doesnt dull dumb ed exactly f
ake feel figure finds gets goes got great guy guys hell help hes hey
hospital hour house huge interested ive job just killer lady lame lat
er leaves like little lives media minute minutes mom oh original plai
n police president real really reporter right secret sharp shes sight
sister spoiler story street stuff suicide talking thats theres theyre
think time told took trailer true trying tv wait wasnt watching whats
white wife working yeah youve

Prediction: 1

Class: 0

actor air attempts bloody bodies capture chance church continue cross
detective diamond die dont dying effective evil film films finish fli
ck forces gas gate gets griffith hand hands happen heart heavy help h
es host ill kill killer know like line major merely need normally par
tner patrick poor power problem pure really satan serial sister sort
soul stand stop supporting tried try type understanding villain weak
welcome women world

Prediction: 1

Class: 0

adult atmosphere bad boy brother camera cast characters christopher c
lose despite didnt director drama experience family film given matter
memorable movie natural nearly necessary older people picture project
really robert sad shooting situation small story subject talent treat
ment unfortunately worth young

Prediction: 1

Class: 0

film forgotten late little long makers money movie night present subt
le time todays tv

Prediction: 0

Class: 1

ago based best big blood close copy ending enjoy fight funny glad got
guess hey hurt im jesus kill killer like looks lot man memory mention
movie scenes screen seen seven shes shown slasher spoil story surely
suspense thats theres times ups video wearing whos wont wood woods wo
rk years

Prediction: 1

Class: 0

annoying beautiful easily escape films finds force george god goes good great hand high island isnt jane johnny mention people period player present priest robert singer song thing tries try used woman worst worthy years

Prediction: 1

Class: 0

action adventure adventures bad camp character characters check crew decided doc elements familiar fan fans feel feeling film good hero heroes im james jones just know long lot major minutes movie movies music number ones promise provided really resulting savage say seeing somewhat spirit star thats theres throw time trying unfortunate way

Prediction: 0

Class: 1

black blood chair course days effects fans film finally gang guts horror idea killing make makers making men murder nasty poor possible presents pretty real seen series sick snuff special thing thrown torture tried water woman

Prediction: 0

Class: 1

acting animals best better die dont entire episode episodes funny good horrible ice just killing know life like movie obviously plot problem really remember right scene scenes season second series shocking suspense think torture turns victims watch women wonderful worst

Prediction: 0

Class: 1

Exercise: Provide explanation as to why these examples were misclassified below.

Type your answer here

The main aim of the model is to predict the sentiment. As each word is taken individually, the model is failing to learn the sentiment of the current word from previous words and is failing to learn from the sentence as a whole.

In []: